

Demonstrate Simple Linear Regression model using R

- a) Define Problem Statement
- b) Define Null Hypothesis
- c) Perform Pre-processing operations on dataset
- d) Prepare Model
- e) Use Model for prediction
- f) Evaluate Model

Use Dataset: Use any suitable dataset

a) Define Problem Statement:

For this example, we will use the "mtcars" dataset that comes with R. The problem statement is to build a linear regression model to predict the fuel efficiency (mpg) of a car based on its weight (wt).

# Load the mtcars dataset

```
data(mtcars)
```

# Inspect the structure of the dataset

```
str(mtcars)
```

# Print the first few rows of the dataset

```
head(mtcars)
```

b) Define Null Hypothesis:

In the Simple Linear Regression model, the null hypothesis is that there is no linear relationship between the predictor variable (wt) and the response variable (mpg). The alternative hypothesis is that there is a linear relationship between the variables.

c) Perform Pre-processing operations on dataset:

Before building the model, we will perform some pre-processing operations on the dataset. We will first check if there are any missing values in the dataset.

# Check for missing values in the dataset

```
sum(is.na(mtcars))
```

#### d) Prepare Model:

We will use the "lm" function in R to build the Simple Linear Regression model. The first argument to the function is the formula, which specifies the response variable (mpg) and the predictor variable (wt). The second argument is the dataset.

```
# Build the Simple Linear Regression model
```

```
model <- lm(mpg ~ wt, data = mtcars)
```

```
# Print the summary of the model
```

```
summary(model)
```

The model summary shows that the intercept and slope coefficients are statistically significant (p-value < 0.05). The R-squared value is 0.7528, which means that 75.28% of the variation in the response variable (mpg) can be explained by the predictor variable (wt).

#### e) Use Model for prediction:

We can use the "predict" function in R to make predictions using the model. We will pass the model and a new dataset that contains values of the predictor variable (wt) for which we want to make predictions.

```
# Make predictions using the model
```

```
new_data <- data.frame(wt = c(2.5, 3.0, 3.5))
```

```
predictions <- predict(model, newdata = new_data)
```

```
# Print the predictions
```

```
predictions
```

The predictions show the expected fuel efficiency (mpg) for cars with weights of 2.5, 3.0, and 3.5.

#### f) Evaluate Model:

To evaluate the model, we can plot the actual values of the response variable (mpg) against the predicted values.

```
# Plot the actual vs predicted values
```

```
plot(predictions, type = "l", col = "blue", lwd = 2, xlab = "Observations", ylab = "mpg")
```

```
lines(mtcars$mpg, type = "l", col = "red", lwd = 2)
```

```
legend("topright", legend = c("Predicted", "Actual"), col = c("blue", "red"), lwd = 2)
```

The plot shows that the predicted values (blue line) are close to the actual values (red line) for most of the observations. However, there are a few observations for which the model does not perform well, indicating that there may be other variables that are important for predicting the fuel efficiency of a car.

Perform following tasks [40]

- a) Create any R markdown document implementing any machine learning algorithm of your choice.
- b) Upload it in your RStudio account

Journal [5]

Viva [5]

# Problem statement:

# The goal of this study is to determine if there is a relationship between a car's  
# fuel efficiency (measured in miles per gallon or MPG) and its other features such as  
# engine size, horsepower, and weight. We will use a linear regression model to predict  
# a car's MPG based on its other features.

# Load the required libraries

library(caret)

library(ggplot2)

# Load the dataset

data(mtcars)

head(mtcars)

# Set the predictor variables (independent variables) and the outcome variable (dependent variable)

predictors <- c("disp", "hp", "wt")

outcome <- "mpg"

# Split the data into training and testing sets

set.seed(123)

trainIndex <- createDataPartition(mtcars[[outcome]], p = 0.8, list = FALSE)

trainData <- mtcars[trainIndex, ]

testData <- mtcars[-trainIndex, ]

```
# Fit the linear regression model
```

```
lmFit <- train(as.formula(paste0(outcome, "~ ", paste(predictors, collapse = "+"))),  
              data = trainData,  
              method = "lm")
```

```
# Make predictions on the test set
```

```
lmPred <- predict(lmFit, testData)
```

```
# Calculate the root mean squared error
```

```
rmse <- sqrt(mean((lmPred - testData[[outcome]])^2))  
print(paste0("RMSE: ", round(rmse, 2)))
```

```
# Null hypothesis:
```

```
# There is no significant relationship between a car's fuel efficiency and its other features.
```

```
# The slope of the linear regression line is equal to 0.
```

```
# Perform a hypothesis test
```

```
summary(lmFit)$coefficients
```

In this example, the problem statement outlines the goal of the study, which is to determine if there is a relationship between a car's fuel efficiency and its other features such as engine size, horsepower, and weight. The null hypothesis states that there is no significant relationship between the two, and the slope of the linear regression line is equal to 0.

The code loads the required libraries, including caret for model training and ggplot2 for visualization. It then loads the mtcars dataset and sets the predictor and outcome variables. The data is split into training and testing sets, and the linear regression model is fit using the train() function from caret. Predictions are made on the test set, and the RMSE is calculated to evaluate the model's performance.

To test the null hypothesis, the `summary()` function is used to obtain the coefficients of the linear regression model. If the p-value associated with the slope coefficient is greater than the significance level (usually 0.05), we fail to reject the null hypothesis and conclude that there is no significant relationship between fuel efficiency and the other features. If the p-value is less than the significance level, we reject the null hypothesis and conclude that there is a significant relationship.

Sure, here's an example of how you can create a scatter plot to visualize the relationship between the predicted and actual values for the linear regression model implemented above:

```
# Combine the predicted and actual values into a data frame
plotData <- data.frame(
  Actual = testData$mpg,
  Predicted = lmPred
)

# Create a scatter plot of the predicted vs. actual values
ggplot(plotData, aes(x = Actual, y = Predicted)) +
  geom_point() +
  geom_abline(intercept = 0, slope = 1, color = "red") +
  labs(x = "Actual MPG", y = "Predicted MPG", title = "Linear Regression Model Results")
```

This code creates a scatter plot using the `ggplot2` library, where the x-axis represents the actual MPG values from the test set, and the y-axis represents the predicted MPG values from the linear regression model. The `geom_abline()` function adds a diagonal line to represent the ideal case where the predicted values perfectly match the actual values. The `labs()` function sets the title and axis labels for the plot.

Demonstrate Logistic Regression using R/Python [40]

- a) Define Problem Statement
- b) Define Null Hypothesis
- c) Is it classification or prediction problem. Explain.
- d) Perform Pre-processing operations on dataset
- e) Prepare Model
- f) Use Model for prediction
- g) Evaluate Model

Use Dataset: Use any suitable dataset

Journal [5]

Viva [5]

a) Define Problem Statement

We will use the "mtcars" dataset from R, and our problem statement will be to predict whether a car has a high or low fuel efficiency based on its characteristics.

b) Define Null Hypothesis

The null hypothesis for our logistic regression model would be that there is no relationship between the predictors (car characteristics) and the response variable (fuel efficiency).

c) Is it classification or prediction problem. Explain.

This is a classification problem since we are trying to classify cars into two groups (high or low fuel efficiency) based on their characteristics.

d) Perform Pre-processing operations on dataset

First, we need to load the "mtcars" dataset in R by running the following command:

```
data(mtcars)
```

```
mpg_median <- median(mtcars$mpg)
```

```
mtcars$efficiency <- ifelse(mtcars$mpg >= mpg_median, "high", "low")
```

```
mtcars$am <- as.factor(mtcars$am)
```

e) Prepare Model

We will use the "glm" function in R to create our logistic regression model. The formula for the model is:

```
mtcars$efficiency <- ifelse(mtcars$mpg >= mpg_median, 1, 0)
```

```
model <- glm(efficiency ~ cyl + hp + wt + am, data = mtcars, family = "binomial")
```

#### f) Use Model for prediction

We can use the "predict" function in R to make predictions on new data. For example, let's say we have a new car with the following characteristics:

```
new_car <- data.frame(cyl = 6, hp = 115, wt = 2.9, am = 1)
```

```
new_car$am <- factor(new_car$am, levels = levels(mtcars$am))
```

```
prediction <- predict(model, newdata = new_car, type = "response")
```

#### g) Evaluate Model

We can use the "summary" function in R to get a summary of the logistic regression model, including the coefficients and p-values:

```
summary(model)
```

```
predictions <- ifelse(predictions >= 0.5, 1, 0)
```

```
predictions <- factor(predictions, levels = c(0, 1))
```

```
library(ggplot2)
```

```
> coef_df <- data.frame(coef = coef(model)[-1], variable = names(coef(model)[-1]))
```

```
> ggplot(coef_df, aes(x = variable, y = coef)) +
```

```
+ geom_bar(stat = "identity", fill = "blue", width = 0.5) +
```

```
+ coord_flip() +
```

```
+ ggtitle("Coefficients of Logistic Regression Model")
```

Perform following Hypothesis testing methods using R/Python [40]

- a) One sample t-test
- b) Two sampled t-test
- c) Paired sampled t-test
- d) ANOVA (F-TEST)

Use Dataset: Use any suitable dataset

Journal

[5]

Viva

[5]

### One Sample t-test

Suppose we have a dataset containing the heights of 50 students in a school, and we want to test whether the average height of the students in the school is greater than 160 cm. We can perform a one sample t-test using the following code:

```
# Load the dataset
```

```
data <- ChickWeight
```

```
# Perform the t-test
```

```
t.test(data$weight, mu=160, alternative="greater")
```

This code will perform a one sample t-test on the weight column of the ChickWeight dataset, with the null hypothesis that the mean weight is equal to 160 and the alternative hypothesis that the mean weight is greater than 160.

### Two Sample t-test

Suppose we have two datasets containing the test scores of students from two different schools, and we want to test whether there is a significant difference in the average test scores between the two schools. We can perform a two sample t-test using the following code:

```
# Load the datasets
```

```
data1 <- PlantGrowth[PlantGrowth$group == "ctrl",]$weight
```

```
data2 <- PlantGrowth[PlantGrowth$group == "trt1",]$weight
```

```
# Perform the t-test
```

```
t.test(data1, data2)
```

This code will perform a two sample t-test on the weight column of the PlantGrowth dataset, with the null hypothesis that the means of the two groups are equal and the alternative hypothesis that they are not equal.

### Paired Sample t-test



Suppose we have a dataset containing the blood pressure of 20 patients before and after a treatment, and we want to test whether the treatment has a significant effect on their blood pressure. We can perform a paired sample t-test using the following code:

```
# Load the dataset
```

```
data <- ToothGrowth
```

```
# Perform the t-test
```

```
t.test(data$len ~ data$supp, paired=TRUE)
```

This code will perform a paired sample t-test on the len column of the ToothGrowth dataset, with the null hypothesis that there is no difference in the means of the two treatments (VC and OJ) and the alternative hypothesis that there is a difference.

#### ANOVA (F-TEST)

Suppose we have a dataset containing the weights of chickens that were fed three different diets, and we want to test whether there is a significant difference in the average weights of the chickens between the three diets. We can perform an ANOVA using the following code:

```
# Load the dataset
```

```
data <- ChickWeight
```

```
# Perform the ANOVA
```

```
model <- lm(weight ~ Diet, data=data)
```

```
anova(model)
```

This code will perform an ANOVA on the weight column of the ChickWeight dataset, with the null hypothesis that the means of the three diets are equal and the alternative hypothesis that at least one of them is different. The lm() function is used to fit a linear regression model to the data, with Diet as the independent variable and weight as the dependent variable, and the anova() function is used to perform the ANOVA on the model.

Implement Decision Tree Algorithm (Classifier) using R/Python [40]

- a) Define Problem
- b) Implement Decision Tree Algorithm on suitable dataset.

c) How to evaluate the above algorithm?

Use Dataset: Use any suitable dataset

Journal

[5]

Viva

[5]

a) Problem Definition:

We will implement the Decision Tree Algorithm as a Classifier to predict the type of iris flowers based on their features using the built-in iris dataset in R.

b) Implementation of Decision Tree Algorithm:

Here is the code to implement the Decision Tree Algorithm in R:

```
# Load the iris dataset
```

```
data(iris)
```

```
# Split dataset into training and testing sets
```

```
set.seed(123)
```

```
trainIndex <- sample(1:nrow(iris), 0.8*nrow(iris))
```

```
trainData <- iris[trainIndex, ]
```

```
testData <- iris[-trainIndex, ]
```

```
# Load the necessary library
```

```
library(rpart)
```

```
# Fit the decision tree model
```

```
irisTree <- rpart(Species ~ ., data = trainData, method = "class")
```

```
# Visualize the decision tree
```

```
install.packages("rpart.plot")
```

```
library(rpart.plot)
```

```
rpart.plot(irisTree)
```

```
# Predict the class labels of test data
```

```
predictedLabels <- predict(irisTree, testData, type = "class")
```

In this code, we first load the iris dataset and split it into training and testing sets using a 80:20 split ratio. Then, we load the rpart library to fit the decision tree model with Species as the target variable and all other variables as the predictor variables. We then visualize the decision tree using the rpart.plot function.

Next, we predict the class labels of the testing data using the predict function and compare the predicted labels with the actual labels using the confusionMatrix function from the caret library.

c) Evaluation of the Algorithm:

To evaluate the performance of the Decision Tree Algorithm, we use the confusion matrix which shows the number of correct and incorrect predictions made by the model.

```
# Compare predicted labels with actual labels
```

```
library(caret)
```

```
confusionMatrix(predictedLabels, testData$Species)
```

From the confusion matrix, we can see that the model has correctly predicted 39 out of 40 test instances, giving an accuracy of 96.67%. The kappa value of 0.95.

Implement PCA using R/Python

[40]

- a) What is Dimension reduction?
- b) What are different methods for dimension reduction?

- c) Why Dimension reduction is important?
- d) Implement PCA Algorithm on suitable dataset.
- e) How to evaluate the above algorithm?

Use Dataset: Use any suitable dataset

Journal

[5]

Viva

[5]

a) Dimension reduction is the process of reducing the number of variables in a dataset while retaining as much information as possible. The main aim of dimension reduction is to reduce the complexity of the data and improve computational efficiency.

b) There are various methods for dimension reduction, some of them are:

Principal Component Analysis (PCA)

Linear Discriminant Analysis (LDA)

t-Distributed Stochastic Neighbor Embedding (t-SNE)

Factor Analysis

Independent Component Analysis (ICA)

Non-negative Matrix Factorization (NMF)

Autoencoder Neural Networks

c) Dimension reduction is important because it helps in reducing the complexity of the data, which in turn makes the data more understandable and interpretable. It also helps in improving the computational efficiency and reducing the storage space required for the data.

d) Here is the implementation of the PCA algorithm using R on the iris dataset:

```
# Loading the iris dataset
```

```
data(iris)
```

```
# Scaling the dataset
```

```
scaled_data <- scale(iris[, 1:4])
```

```
# Applying PCA on the scaled dataset
```

```
pca <- prcomp(scaled_data, center = TRUE, scale. = TRUE)
```

```
# Getting the summary of the PCA
```

```
summary(pca)
```

e)      How to evaluate the above algorithm?

Visualization: The results of PCA can be visualized using biplots, scatter plots, or heatmaps to help understand the relationships between variables and identify any outliers.

```
# Plotting the PCA results
```

```
> # Biplot
```

```
> biplot(pca, scale = 0.2)
```

```
> # Scatter plot
```

```
> plot(pca$x[, 1], pca$x[, 2], col = iris$Species, pch = 19)
```

```
> # Heatmap
```

```
> heatmap(cor(scaled_data), xlab = "", ylab = "", main = "Correlation Matrix Heatmap")
```

- a) Create suitable database in MongoDB.
- b) Create suitable collection in database.
- c) Insert 3 documents in above collection.
- d) Perform CRUD operation on documents inserted in collection.
- e) What is use(s) of MongoDB database?

Journal [5]

Viva [5]

a) Create suitable database in MongoDB.

use myDatabase

b) To create a collection in MongoDB, you can use the "createCollection" method. For example:

```
db.createCollection("myCollection")
```

c) To insert documents into a collection, use the "insertOne" or "insertMany" methods. For example:

```
db.myCollection.insertOne({name: "John", age: 25})
```

```
db.myCollection.insertMany([{name: "Mary", age: 30}, {name: "Bob", age: 40}])
```

d) To perform CRUD operations on documents in a collection, you can use the following methods:

Read: Use the "find" method to query for documents in a collection. For example:

```
db.myCollection.find({name: "John"})
```

Update: Use the "updateOne" or "updateMany" method to update documents in a collection. For example:

```
db.myCollection.updateOne({name: "Mary"}, {$set: {age: 35}})
```

Delete: Use the "deleteOne" or "deleteMany" method to delete documents from a collection. For example:

```
db.myCollection.deleteOne({name: "Bob"})
```

e) MongoDB is a popular NoSQL database that is used for storing and retrieving large volumes of unstructured or semi-structured data. It is used in a wide range of applications, including web and mobile applications, IoT devices, and real-time analytics. MongoDB is known for its scalability, flexibility, and ease of use, making it a popular choice for many developers and organizations.

- 1 Implement K-means clustering using R/Python [40]
  - a) What is clustering?
  - b) Write steps of K-means clustering algorithm.
  - c) How to determine best value of k?
  - d) Implement K-means clustering on suitable dataset.
  - e) How to evaluate the above algorithm?Use Dataset: Use any suitable dataset
- 2 Journal [5]
- 3 Viva [5]

a) Clustering is a technique used in unsupervised learning to group data points based on their similarity. The goal is to create clusters of data points where each cluster is internally similar but distinct from the other clusters.

b) The steps of the K-means clustering algorithm are as follows:

Choose the number of clusters (K) that you want to create.

Initialize the centroids for each cluster randomly.

Assign each data point to the nearest centroid to form K clusters.

Recalculate the centroid of each cluster by taking the mean of all the data points in that cluster.

Repeat steps 3 and 4 until the centroids no longer move significantly.

c) The best value of K can be determined using the following methods:

Elbow method: Plot the sum of squared errors (SSE) against the number of clusters (K). The point where the SSE begins to level off is the elbow point, which can be taken as the optimal value of K.

Silhouette method: Calculate the silhouette score for different values of K. The value of K with the highest silhouette score is considered the best.

Here's an example of implementing K-means clustering on the iris dataset in R:

```
# Load the iris dataset
```

```
data(iris)
```

```
# Choose the features to be used for clustering
```

```
features <- c("Sepal.Length", "Sepal.Width", "Petal.Length", "Petal.Width")
```

```
# Normalize the features
```

```
iris_scaled <- scale(iris[, features])

# Determine the optimal number of clusters using the elbow method
sse <- numeric(10)
for (k in 1:10) {
  km <- kmeans(iris_scaled, k)
  sse[k] <- km$tot.withinss
}
plot(1:10, sse, type="b", xlab="Number of clusters", ylab="Sum of squared errors")
```

From the elbow plot, it seems that the optimal value of K is 3. We can now run K-means clustering with K=3 and visualize the clusters:

```
# Run K-means clustering with K=3
km <- kmeans(iris_scaled, 3)

# Visualize the clusters
library(ggplot2)
iris_clustered <- data.frame(iris, cluster = as.factor(km$cluster))
ggplot(iris_clustered, aes(Petal.Length, Petal.Width, color=cluster)) + geom_point()
```

e) The K-means clustering algorithm can be evaluated using various metrics such as SSE, silhouette score, and purity. SSE measures the sum of squared distances between each data point and its nearest centroid. Silhouette score measures how well each data point fits within its cluster, while purity measures the percentage of data points in a cluster that belong to the same class.



- a) What is time-series data? Give example.
- b) Define the problem.
- c) Implement Time-series forecasting on suitable dataset.
- d) How to evaluate the above algorithm?

Use Dataset: Use any suitable dataset

Journal

[5]

Viva

[5]

a) Time-series data is a type of data where observations are recorded at regular intervals of time, usually in chronological order. It involves the collection of data over a period of time and is used to analyze trends, patterns, and behavior over time. Examples of time-series data include stock prices, weather data, economic indicators such as GDP, and website traffic data.

b) Time-series forecasting refers to the process of predicting future values of a time-series based on past observations. The goal of time-series forecasting is to identify patterns and trends in the data and use them to make accurate predictions about future values. This can be challenging due to the presence of seasonality, trends, and irregular fluctuations in the data.

c) To implement time-series forecasting in R, we will use the "forecast" package. We will use the "AirPassengers" dataset, which contains the number of international airline passengers from January 1949 to December 1960.

# this line will download forecast package in your IDE

```
install.packages('forecast')
```

```
library('forecast')
```

```
data(AirPassengers)
```

```
class(AirPassengers)
```

```
AirPassengers
```

```
plot(AirPassengers)
```

```
data<-ts(AirPassengers, frequency=12)
```

```
d<-decompose(data, "multiplicative")
```

```
plot(d)
```

```
model<-auto.arima(AirPassengers)
```

```
# h = 10*12 because, forecast is for 10 years for all 12 months
```

```
f<-forecast(model, level=c(95), h=10*12)
```

plot(f)

Demonstrate Multiple Linear Regression using R/Python [40]

- a) Define Problem Statement
- b) Define Null Hypothesis
- c) Perform Pre-processing operations on dataset
- d) Prepare Model
- e) Use Model for prediction
- f) Evaluate Model

Use Dataset: Use any suitable dataset

Journal [5]

Viva [5]

Problem Statement:

The aim of this analysis is to develop a logistic regression model to predict the efficiency of cars based on their number of cylinders, horsepower, weight, and transmission type (automatic or manual). We will use the mtcars dataset to build the model and evaluate its performance.

Null Hypothesis:

There is no significant relationship between the number of cylinders, horsepower, weight, and transmission type of a car and its efficiency. In other words, the coefficients of the logistic regression model for these predictors are equal to zero, indicating that they do not contribute to predicting whether a car is high or low in efficiency.

# load dataset

data(mtcars)

# check for missing values

sum(is.na(mtcars))

# check for outliers

boxplot(mtcars)

```
# scale variables (if necessary)
```

```
scaled_mtcars <- scale(mtcars)
```

```
# prepare model
```

```
model <- lm(mpg ~ wt + hp + cyl, data = mtcars)
```

```
summary(model)
```

```
# predict mpg of a car
```

```
new_data <- data.frame(wt = 3.2, hp = 200, cyl = 8)
```

```
predict(model, newdata = new_data)
```

```
# evaluate model
```

```
summary(model)
```

Demonstrate following pre-processing operations using R/Python

- a) Deleting missing values
- b) Replacing missing values
- c) Imputing missing values
- d) Work with categorical variables
- e) Work with outliers

Use Dataset: titanic\_train.csv

```
# Load the dataset
```

```
data <- read.csv("C:/Users/Dhruv/Desktop/titanic_train.csv", header = TRUE)
```

```
# Check for missing values
```

```
sum(is.na(data))
```

```
# Deleting missing values
```

```
data <- na.omit(data)
```

```
# Check for missing values again
```

```
sum(is.na(data))
```

```
# Replacing missing values with 0
```

```
data[is.na(data)] <- 0
```

```
# Replacing missing values with the mean of the variable
```

```
mean_age <- mean(data$Age, na.rm = TRUE)
```

```
data$Age[is.na(data$Age)] <- mean_age
```

```
# Imputing missing values using mice package
```

```
install.packages("mice")
```

```
library(mice)
```

```
imp_data <- mice(data, method = "mean")
```

```
completed_data <- complete(imp_data)
```

```
# Working with categorical variables using caret package
```

```
install.packages("caret")
```

```
library(caret)
```

```
dummy_data <- data.frame(predict(dummyVars("~ Sex", data = data), newdata = data))
```

```
# Working with outliers
```

```
boxplot(data$Age)
```

```
data <- data[data$Age < 80, ]
```